

1 IMPLEMENTATION OF BOX STENCIL KERNEL

Algorithm 1 Box stencil kernel in *TCstencil*

```

1: Input: mesh A, para_box1, para_box2, para_box3
2: Output: the updated mesh C'
3: Shared Memory: sinput, soutput1,soutput2,soutput3
4: Fragment: a_frag,b_frag,c_frag
5: p=[para_box1, para_box2, para_box3]
6: sout=[soutput1,soutput2,soutput3]
7: load A' of A to sinput
8: for i from 0 to 2 do
9:   wmma::fill_fragment(c_frag, 0.0f)
10:  wmma::load_matrix_sync(a_frag, p[i], 16)
11:  wmma::load_matrix_sync(b_frag, sinput, 16)
12:  wmma::mma_sync(c_frag, a_frag, b_frag, c_frag)
13:  wmma::store_matrix_sync(sout[i], c_frag, 16, wmma::mem_row_major)
14: calculate result by soutput1,soutput2,soutput3 and store it to global memory

```

2 ROOFLINE ANALYSIS

We adopt the roofline model [2] to analyze the effectiveness of stencil computation by Artemis, *TCstencil*, and *TC-w/o tc*. **Since Nsight profiler do not support collect FLOP in half, so we give the theoretical value.** Table 1 shows the the maximum operational intensity $I_{max} = \pi/\beta$ (FLOP/Byte) in FP16 precision of A100 and V100, where π (TFLOPS) is the maximum attainable performance the β (GB/s) is maximum memory bandwidth. We are only considering the multiplication-and-add for stencil as FLOP and the access to global memory. For updating the inner points in 16×16 mesh with 9pt stencil, the Arithmetic Intensity (AI, FLOP/Byte) of Artemis/*TC-w/o tc*/*TCstencil* is 0.50/2.53/5.33. As for 25pt stencil, the AI of Artemis/*TC-w/o tc*/*TCstencil* is 0.50/7.03/6. The computation details are in the appendix. All the above implementations are memory bound since their AI are lower than I_{max} in Table 1.

Table 1: The performance information of A100 and V100 in FP16 precision [1].

Item	A100 FMA	A100 TC	V100 FMA	V100 TC
π	78	312	31.4	125
β	1555		900	
I_{max}	50.1	200.64	34.88	128.8

3 APPENDIX

The computation detail:

- Artemis 9pt: 9 FLOP/(9*2) Byte. update every point, read 9 points in half and perform 9 FLOP.
- Artemis 25pt: 25 FLOP/(25*2) Byte.
- *TC-w/o tc* 9pt: (12*12*9) FLOP / (16*16*2) Byte. update every 12*12 points, load 16*16 points in global memory, perform 12*12*9 FLOP.
- *TC-w/o tc* 25pt: (12*12*25) FLOP / (16*16*2) Byte.
- *TCstencil* 9pt: (16*16*16*2) FLOP/ (3*16*16)*2 Byte. update every 12*12 point, load 16*16 points and 2*16*16 elements in parameter matrices. perform twice Tensor Core MMA, each MMA perform 16*16*16 FLOP.
- *TCstencil* 25pt: (16*16*16*3) FLOP/ (4*16*16)*2 Byte.

REFERENCES

- [1] NVIDIA. 2021-1-30. NVIDIA Turing Architecture Whitepaper. online. <https://www.nvidia.com/en-us/data-center/a100/> Accessed March 30, 2021.
- [2] Samuel Williams, Andrew Waterman, and David Patterson. 2009. Roofline: an insightful visual performance model for multicore architectures. *Commun. ACM* 52, 4 (2009), 65–76.